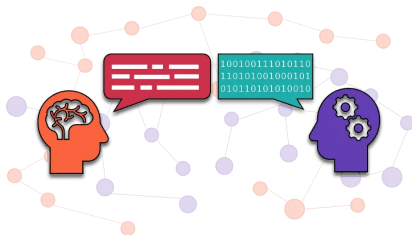


# Large Language Models

**Konstantinos Skianis**, Assistant Professor  
Department of Computer Science & Engineering  
University of Ioannina, Greece



**Makeathon, UniAI**

Friday 19/04/2024

## 1 Introduction

## 2 Applications and Use Cases

## 3 Prompt Engineering

## 4 Retrieval Augmented Generation

## 5 Summary

# Introduction

## Large Language Models (LLMs)

Based on the Transformer architecture, which has become the foundation for various state-of-the-art natural language processing (NLP) models

### Transformer:

- Attention is All You Need ([Vaswani et al., 2017](#))
- a groundbreaking neural network architecture designed for NLP tasks
- relies on the self-attention mechanism to process and generate sequences
- highly efficient and scalable compared to traditional recurrent neural networks (RNNs) and long short-term memory (LSTM) models

# Variants of Transformer Architecture

Various LLMs are built on top of the Transformer architecture with slight modifications or adaptations. Some popular variants include:

- **GPT**: The Generative Pre-trained Transformer (GPT) is an autoregressive model that utilizes only the decoder part of the Transformer architecture to generate text.
- **BERT**: Bidirectional Encoder Representations from Transformers (BERT) is based on the encoder part of the Transformer architecture and is pre-trained using masked language modeling and next sentence prediction tasks.
- **T5**: The Text-to-Text Transfer Transformer (T5) adapts the original Transformer architecture to a unified text-to-text format, enabling it to be used for various NLP tasks with minimal task-specific modifications.

1 Introduction

2 Applications and Use Cases

3 Prompt Engineering

4 Retrieval Augmented Generation

5 Summary

# Applications and Use Cases of LLMs

- **Sentiment Analysis:** can be used or fine-tuned to analyze the sentiment expressed in a piece of text, determining whether it is positive, negative, or neutral. → gauge customer satisfaction, monitor social media sentiment, or analyze product reviews
- **Question-Answering:** build question-answering systems that provide precise answers to user queries → general or domain-specific
- **Text Summarization:** generate a concise summary of a given input text. This can be useful for creating executive summaries of articles, reports, news stories, or research studies making it easier for readers to quickly grasp the main points
- **Machine Translation:** can be adapted for machine translation tasks, where the goal is to translate text from one language to another. By fine-tuning on parallel corpora, models like T5 have demonstrated state-of-the-art performance in many language pairs

## Applications and Use Cases of LLMs (2)

- **Named Entity Recognition (NER):** can be used to identify and classify named entities (e.g., people, organizations, locations) within a text  
→ information extraction, content recommendation, and semantic search
- **Text Classification:** classify text into various categories  
→ topic classification, spam detection, or document tagging content filtering and recommendation systems
- **Text Generation:** LLMs, especially autoregressive models like GPT, can be used for text generation tasks  
→ generate coherent and contextually relevant text based on a given input, which can be useful for tasks like chatbot development, content creation, or creative writing prompts
- **Few-Shot Learning:** can adapt to new tasks with minimal task-specific examples → enable more efficient transfer learning and reduce the need for extensive fine-tuning

- 1 Introduction
- 2 Applications and Use Cases
- 3 Prompt Engineering**
- 4 Retrieval Augmented Generation
- 5 Summary



# Prompt Engineering

- (1) Explicitness: specify the format you want:
  - Less explicit prompt: Translate the following English text to French:  
"Hello, how are you?"
  - More explicit prompt: Translate the following English text to French and provide the translation in quotes: "Hello, how are you?"
- (2) In-context examples: provide examples within the prompt to guide the model towards generating responses in the desired format or style:
  - Without example: Write a haiku about nature.
  - With example: Write a haiku about nature, following the 5-7-5 syllable pattern. For example: "Gentle morning rain, / Caressing the earth below, / Nature's symphony."
- (3) Limit response length: specify a maximum response length to prevent excessively long or verbose answers.
  - Without length limit: Explain the greenhouse effect.
  - With length limit: Explain the greenhouse effect in less than 3 sentences
- (4) Requesting rationales or step-by-step explanations:
  - without rationale: What is the square root of 16?
  - with rationale: What is the square root of 16, and explain why that is the correct answer

# Chain-of-Thought

## Standard Prompting

### Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

### Model Output

A: The answer is 27. ❌

## Chain of Thought Prompting

### Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls.  $5 + 6 = 11$ . The answer is 11.

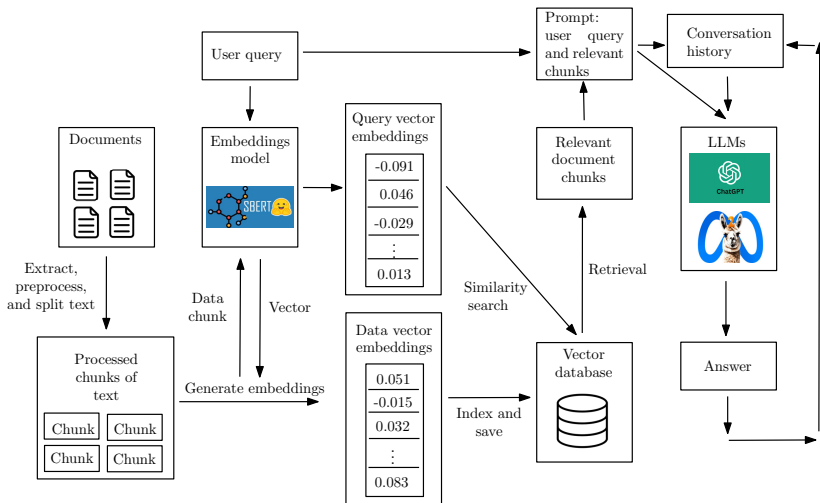
Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

### Model Output

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had  $23 - 20 = 3$ . They bought 6 more apples, so they have  $3 + 6 = 9$ . The answer is 9. ✅

- 1 Introduction
- 2 Applications and Use Cases
- 3 Prompt Engineering
- 4 Retrieval Augmented Generation**
- 5 Summary

# Retrieval Augmented Generation (RAG) Pipeline



A full schematic illustration, based on: <https://ubuntu.com/blog/llms-explained>

# Why use RAG?

## Benefits

- **Consistency:** more likely to get the same answer from the same question
- **Increased factual accuracy:** LLM's responses are based on the provided information → model is less likely to hallucinate or mislead
- **Decreased cost:** building a RAG pipeline is less expensive than fine-tuning. Update only the database instead of training a new model
- **Currency:** ensure the LLM's responses are based on up-to-date data
- **Easily verifiable source:** access to the source for cross-checking
- **Strong scalability:** millions of documents

## Cons

- **Complexity:** chunking, quality of search, low-quality data
- **Context length limitation:** ChatGPT (gpt-3.5-turbo) has a maximum context length of 15k tokens
- **Limited creativity:** losing the point of generative AI

- 1 Introduction
- 2 Applications and Use Cases
- 3 Prompt Engineering
- 4 Retrieval Augmented Generation
- 5 Summary**

# Summary

## We talked about:

- LLMs and applications
- Prompt engineering
- Retrieval Augmented Generation

Link: [https://github.com/cohere-ai/notebooks/blob/main/notebooks/Vanilla\\_RAG.ipynb](https://github.com/cohere-ai/notebooks/blob/main/notebooks/Vanilla_RAG.ipynb)

Coding time!